

## Task 2. Use Sqoop command to ingest the data from RDS into the HBase Table

Logged into EMR cluster

Executed the following command to install MySQL connector jar file

wget <https://de-mysql-connectors3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz>

```
[hadoop@ip-172-31-47-204 ~]$  
[hadoop@ip-172-31-47-204 ~]$  
[hadoop@ip-172-31-47-204 ~]$  
[hadoop@ip-172-31-47-204 ~]$ wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz  
--2024-02-29 08:36:40-- https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz  
Resolving de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)... 52.217.234.169, 52.217.83.228, 52.217.138.49, ...  
Connecting to de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)|52.217.234.169|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 4079310 (3.9M) [application/x-gzip]  
Saving to: 'mysql-connector-java-8.0.25.tar.gz'  
  
100%[=====>] 4,079,310 --K/s in 0.1s  
2024-02-29 08:36:40 (33.6 MB/s) - 'mysql-connector-java-8.0.25.tar.gz' saved [4079310/4079310]
```

Extracted MySQL connector tar file:

tar -xvf mysql-connector-java-8.0.25.tar.gz

```
[hadoop@ip-172-31-47-204 ~]$ tar -xvf mysql-connector-java-8.0.25.tar.gz  
mysql-connector-java-8.0.25/  
mysql-connector-java-8.0.25/src/  
mysql-connector-java-8.0.25/src/build/  
mysql-connector-java-8.0.25/src/build/java/  
mysql-connector-java-8.0.25/src/build/java/documentation/  
mysql-connector-java-8.0.25/src/build/java/instrumentation/  
mysql-connector-java-8.0.25/src/build/misc/  
mysql-connector-java-8.0.25/src/build/misc/debian.in/  
mysql-connector-java-8.0.25/src/build/misc/debian.in/source/  
mysql-connector-java-8.0.25/src/demo/  
mysql-connector-java-8.0.25/src/demo/java/  
mysql-connector-java-8.0.25/src/demo/java/demo/  
mysql-connector-java-8.0.25/src/demo/java/demo/x/  
mysql-connector-java-8.0.25/src/demo/java/demo/x/devapi/  
mysql-connector-java-8.0.25/src/generated/  
mysql-connector-java-8.0.25/src/generated/java/  
mysql-connector-java-8.0.25/src/generated/java/com/  
mysql-connector-java-8.0.25/src/generated/java/com/mysql/  
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/  
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/x/  
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/x/protobuf/  
mysql-connector-java-8.0.25/src/legacy/  
mysql-connector-java-8.0.25/src/legacy/java/  
mysql-connector-java-8.0.25/src/legacy/java/com/  
mysql-connector-java-8.0.25/src/legacy/java/com/mysql/  
mysql-connector-java-8.0.25/src/legacy/java/com/mysql/ids/  
mysql-connector-java-8.0.25/src/legacy/java/com/mysql/ids/
```

Changed the directory to the MySQL Connector directory and copied to the Sqoop library as follows:

- cd mysql-connector-java-8.0.25/
- sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/

```
hadoop@ip-172-31-47-204:~/mysql-connector-java-8.0.25  
[hadoop@ip-172-31-47-204 ~]$  
[hadoop@ip-172-31-47-204 ~]$ cd mysql-connector-java-8.0.25/  
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$  
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$  
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/  
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$  
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$ mysql_secure_installation  
  
NOTE: RUNNING ALL PARTS OF THIS SCRIPT IS RECOMMENDED FOR ALL MariaDB  
SERVERS IN PRODUCTION USE! PLEASE READ EACH STEP CAREFULLY!  
  
In order to log into MariaDB to secure it, we'll need the current  
password for the root user. If you've just installed MariaDB, and  
you haven't set the root password yet, the password will be blank,  
so you should just press enter here.  
  
Enter current password for root (enter for none):  
OK, successfully used password, moving on...  
  
Setting the root password ensures that nobody can log into the MariaDB  
root user without the proper authorisation.  
  
Set root password? [Y/n] Y  
New password:  
Re-enter new password:  
Password updated successfully!  
Reloading privilege tables..  
... Success!  
  
By default, a MariaDB installation has an anonymous user, allowing anyone  
to log into MariaDB without having to have a user account created for  
them. This is intended only for testing, and to make the installation  
go a bit smoother. You should remove them before moving into a  
production environment.
```

## Installed & setup MySQL Connector on EMR cluster:

```
hadoop@ip-172-31-47-204:~/mysql-connector-java-8.0.25
- Dropping test database...
... Success!
- Removing privileges on test database...
... Success!

Reloading the privilege tables will ensure that all changes made so far
will take effect immediately.

Reload privilege tables now? [Y/n] Y
... Success!

Cleaning up...

All done! If you've completed all of the above steps, your MariaDB
installation should now be secure.

Thanks for using MariaDB!
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$ mysql -u root -p
Enter password:
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MariaDB connection id is 47
Server version: 5.5.68-MariaDB MariaDB Server

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [(none)]> GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by 'admin123' WITH GR
ANT
-> OPTION;
Query OK, 0 rows affected (0.00 sec)

MariaDB [(none)]> flush privileges;
Query OK, 0 rows affected (0.00 sec)
```

## Accessed HBase shell on EMR using PuTTY

- Switched to ‘root’ user using the **sudo -i** command
- Access the HBase shell by using the “hbase shell” command

```
root@ip-172-31-47-204:~
MariaDB [(none)]> exit;
Bye
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$
[hadoop@ip-172-31-47-204 mysql-connector-java-8.0.25]$ sudo -i

EEEEEEEEEEEEEEEEEEEE MMMMMMMM          MMMMMMMM RRRRRRRRRRRRRRRR
E::::::::::::::::::::E M::::::::M          M::::::::M R::::::::::::R
EE::::::::EEEEEEEEEE:E M::::::::M          M::::::::M R::::RRRRRRR::::R
  E:::E          EEEEE M::::::::M          M::::::::M RR:::R          R:::R
  E:::E          M:::M::M M:::M::M          M:::M::M R:::R          R:::R
  E:::EEEEEEEEEE M:::M M:::M M:::M M:::M          R::RRRRRRR::::R
  E:::::::::::::::E M:::M M:::M::M M:::M          R:::::::::RR
  E:::EEEEEEEEEE M:::M M:::M M:::M          R::RRRRRRR::::R
  E:::E          M:::M M:::M M:::M          R:::R          R:::R
  E:::E          EEEEE M:::M          MMM          M:::M          R:::R          R:::R
EE::::::::EEEEEEEE::E M:::M          M:::M          R:::R          R:::R
E::::::::::::::::::::E M:::M          M:::M          RR:::R          R:::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM          MMMMMMMM RRRRRRR          RRRRRR

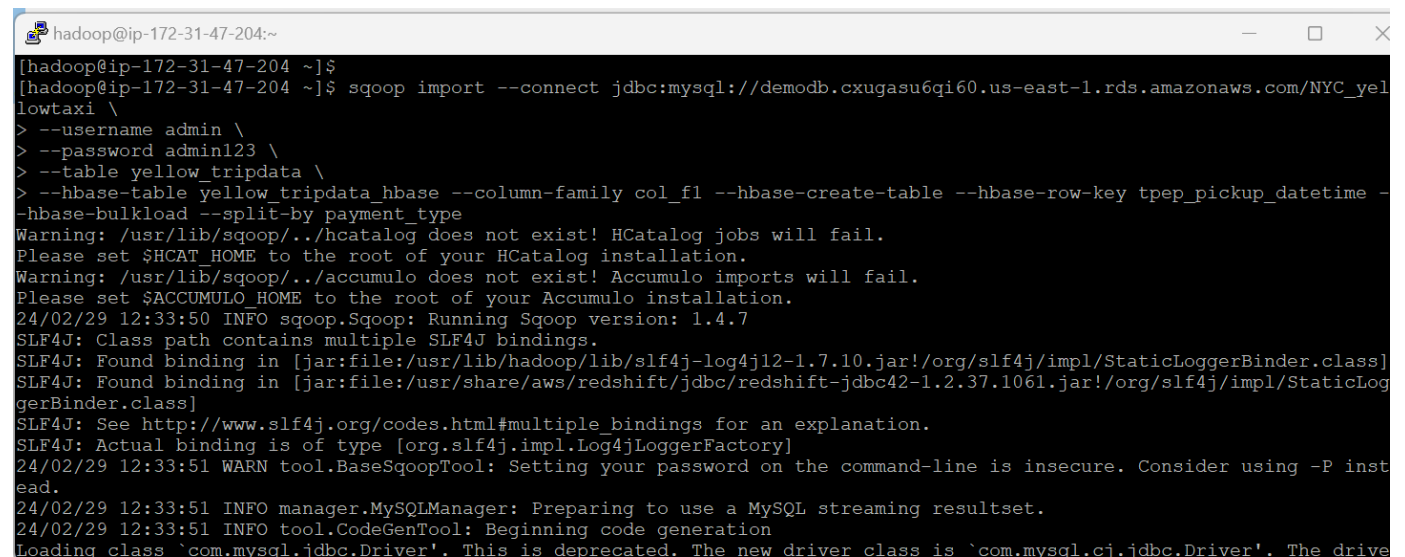
[root@ip-172-31-47-204 ~]# hbase shell
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
Version 1.4.13, rUnknown, Fri Apr 17 15:18:24 UTC 2020
hbase(main):001:0> █
```

Ingested data from RDS to HBase table using following command:

```
sqoop import --connect jdbc:mysql://demodb.cxugasu6qi60.us-east-1.rds.amazonaws.com/NYC_yellowtaxi \  
--username admin \  
--password admin123 \  
--table yellow_tripdata \  
--hbase-table yellow_tripdata_hbase \  
--column-family col_fl --hbase-create-table --hbase-row-key tpep_pickup_datetime --hbase-bulkload  
--split-by payment_type
```

This command transfers data from RDS table “yellow\_tripdata” to an HBase table “yellow\_tripdata\_HBase”.

- connect: JDBC connection string for the RDS(MySQL) database
- username: username to use to connect to the Database
- password: password to use to connect to the Database
- table: name of the source (MySQL) table
- hbase-table: name of the target (HBase) table
- column-family: name of the column family in HBase table
- hbase-create-table: creates HBase table (in case not exists)
- hbase-row-key: column/s of source table to be used as key in HBase table
- hbase-bulkload: Enables bulk loading
- split-by: column used to create the split while importing the data into the cluster



```
hadoop@ip-172-31-47-204:~  
[hadoop@ip-172-31-47-204 ~]$  
[hadoop@ip-172-31-47-204 ~]$ sqoop import --connect jdbc:mysql://demodb.cxugasu6qi60.us-east-1.rds.amazonaws.com/NYC_yel  
lowtaxi \  
> --username admin \  
> --password admin123 \  
> --table yellow_tripdata \  
> --hbase-table yellow_tripdata_hbase --column-family col_fl --hbase-create-table --hbase-row-key tpep_pickup_datetime -  
hbase-bulkload --split-by payment_type  
Warning: /usr/lib/sqoop/./hcatalog does not exist! HCatalog jobs will fail.  
Please set $HCAT_HOME to the root of your HCatalog installation.  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
24/02/29 12:33:50 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7  
SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/usr/lib/hadoop/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/usr/share/aws/redshift/jdbc/redshift-jdbc42-1.2.37.1061.jar!/org/slf4j/impl/StaticLog  
gerBinder.class]  
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.  
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]  
24/02/29 12:33:51 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P inst  
ead.  
24/02/29 12:33:51 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
24/02/29 12:33:51 INFO tool.CodeGenTool: Beginning code generation  
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The drive
```

Checked the HBase table and its statistics created by Sqoop command in HBase shell:

```
hadoop@ip-172-31-47-204:~
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
Version 1.4.13, rUnknown, Fri Apr 17 15:18:24 UTC 2020

hbase(main):001:0> list
TABLE
yellow_tripdata_hbase
1 row(s) in 0.5960 seconds

=> ["yellow_tripdata_hbase"]
hbase(main):002:0> describe yellow_tripdata_hbase
NameError: undefined local variable or method `yellow_tripdata_hbase' for #<Object:0x149aa7b2>

hbase(main):003:0> describe 'yellow_tripdata_hbase'
Table yellow_tripdata_hbase is ENABLED
yellow_tripdata_hbase
COLUMN FAMILIES DESCRIPTION
{NAME => 'col_f1', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
1 row(s) in 0.2640 seconds
```

\*\*\*\*\*

**Task 3.** Bulk import data from next two files in the dataset on your EMR cluster to your HBase Table using the relevant codes.

**Note:** For the above task 3, you just need to import data from the subsequent 2 csv files (i.e. yellow\_tripdata\_2017-03.csv & yellow\_tripdata\_2017-04.csv) on your EMR cluster.

Downloaded the necessary data on local file system.

wget https://nyc-tlc-upgrad.s3.amazonaws.com/yellow\_tripdata\_2017-03.csv

wget https://nyc-tlc-upgrad.s3.amazonaws.com/yellow\_tripdata\_2017-04.csv

```
[hadoop@ip-172-31-47-204 ~]$ wget https://nyc-tlc-upgrad.s3.amazonaws.com/yellow_tripdata_2017-03.csv
--2024-02-29 15:07:45-- https://nyc-tlc-upgrad.s3.amazonaws.com/yellow_tripdata_2017-03.csv
Resolving nyc-tlc-upgrad.s3.amazonaws.com (nyc-tlc-upgrad.s3.amazonaws.com)... 54.231.204.225, 3.5.7.189, 52.217.74.20, ...
Connecting to nyc-tlc-upgrad.s3.amazonaws.com (nyc-tlc-upgrad.s3.amazonaws.com)|54.231.204.225|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 969809025 (925M) [text/csv]
Saving to: 'yellow_tripdata_2017-03.csv'

100%[=====>] 969,809,025 20.6MB/s in 42s

2024-02-29 15:08:26 (22.2 MB/s) - 'yellow_tripdata_2017-03.csv' saved [969809025/969809025]

[hadoop@ip-172-31-47-204 ~]$ wget https://nyc-tlc-upgrad.s3.amazonaws.com/yellow_tripdata_2017-04.csv
--2024-02-29 15:10:13-- https://nyc-tlc-upgrad.s3.amazonaws.com/yellow_tripdata_2017-04.csv
Resolving nyc-tlc-upgrad.s3.amazonaws.com (nyc-tlc-upgrad.s3.amazonaws.com)... 16.182.71.57, 52.216.56.169, 16.182.70.113, ...
Connecting to nyc-tlc-upgrad.s3.amazonaws.com (nyc-tlc-upgrad.s3.amazonaws.com)|16.182.71.57|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 946349441 (903M) [text/csv]
Saving to: 'yellow_tripdata_2017-04.csv'

100%[=====>] 946,349,441 24.0MB/s in 41s

2024-02-29 15:10:54 (22.1 MB/s) - 'yellow_tripdata_2017-04.csv' saved [946349441/946349441]

[hadoop@ip-172-31-47-204 ~]$
```

Setting up the environment to install HappyBase (Python-based HBase API) for executing the batch script as follows:

Installing gcc:

yum install gcc

```
root@ip-172-31-47-204:~  
-rw-r--r-- 1 root root 51868 Feb 29 14:10 yellow_tripdata.java  
-rwxr-x--- 1 root root 1022 Feb 29 17:17 mrtask_a.py  
-rwxr-x--- 1 root root 1821 Feb 29 17:17 Bulk_Import.py  
-rwxr-x--- 1 root root 1584 Feb 29 17:17 mrtask_f.py  
-rwxr-x--- 1 root root 1360 Feb 29 17:17 mrtask_e.py  
-rwxr-x--- 1 root root 1834 Feb 29 17:17 mrtask_d.py  
-rwxr-x--- 1 root root 1389 Feb 29 17:17 mrtask_c.py  
-rwxr-x--- 1 root root 1038 Feb 29 17:17 mrtask_b.py  
-rw-r--r-- 1 root root 969809025 Feb 29 17:18 yellow_tripdata_2017-03.csv  
-rw-r--r-- 1 root root 946349441 Feb 29 17:19 yellow_tripdata_2017-04.csv  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]# yum install gcc  
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd  
amzn2-core | 3.6 kB 00:00:00  
amzn2extra-corretto8 | 2.9 kB 00:00:00  
amzn2extra-docker | 2.9 kB 00:00:00  
emr-apps | 2.5 kB 00:00:00  
emr-platform | 2.5 kB 00:00:00  
14 packages excluded due to repository priority protections  
Resolving Dependencies  
--> Running transaction check
```

Gcc installation completed

```
root@ip-172-31-47-204:~  
Verifying : libmpx-7.3.1-6.amzn2.0.4.x86_64 17/26  
Verifying : libitm-7.3.1-6.amzn2.0.4.x86_64 18/26  
Verifying : libstdc++-7.3.1-6.amzn2.0.4.x86_64 19/26  
Verifying : gcc-gfortran-7.3.1-6.amzn2.0.4.x86_64 20/26  
Verifying : libcilkrts-7.3.1-6.amzn2.0.4.x86_64 21/26  
Verifying : gcc-c++-7.3.1-6.amzn2.0.4.x86_64 22/26  
Verifying : libgcc-7.3.1-6.amzn2.0.4.x86_64 23/26  
Verifying : libatomic-7.3.1-6.amzn2.0.4.x86_64 24/26  
Verifying : libsanitizer-7.3.1-6.amzn2.0.4.x86_64 25/26  
Verifying : cpp-7.3.1-6.amzn2.0.4.x86_64 26/26  
Updated:  
gcc.x86_64 0:7.3.1-17.amzn2  
Dependency Updated:  
cpp.x86_64 0:7.3.1-17.amzn2 gcc-c++.x86_64 0:7.3.1-17.amzn2  
gcc-gfortran.x86_64 0:7.3.1-17.amzn2 libatomic.x86_64 0:7.3.1-17.amzn2  
libcilkrts.x86_64 0:7.3.1-17.amzn2 libgcc.x86_64 0:7.3.1-17.amzn2  
libgomp.x86_64 0:7.3.1-17.amzn2 libitm.x86_64 0:7.3.1-17.amzn2  
libmpx.x86_64 0:7.3.1-17.amzn2 libquadmath.x86_64 0:7.3.1-17.amzn2  
libsanitizer.x86_64 0:7.3.1-17.amzn2 libstdc++.x86_64 0:7.3.1-17.amzn2  
Complete!
```

Installing HappyBase package:

sudo yum install python3-devel

```
root@ip-172-31-47-204:~  
Dependency Updated:  
cpp.x86_64 0:7.3.1-17.amzn2 gcc-c++.x86_64 0:7.3.1-17.amzn2  
gcc-gfortran.x86_64 0:7.3.1-17.amzn2 libatomic.x86_64 0:7.3.1-17.amzn2  
libcilkrts.x86_64 0:7.3.1-17.amzn2 libgcc.x86_64 0:7.3.1-17.amzn2  
libgomp.x86_64 0:7.3.1-17.amzn2 libitm.x86_64 0:7.3.1-17.amzn2  
libmpx.x86_64 0:7.3.1-17.amzn2 libquadmath.x86_64 0:7.3.1-17.amzn2  
libsanitizer.x86_64 0:7.3.1-17.amzn2 libstdc++.x86_64 0:7.3.1-17.amzn2  
Complete!  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]# sudo yum install python3-devel  
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd  
14 packages excluded due to repository priority protections  
Resolving Dependencies  
--> Running transaction check  
--> Package python3-devel.x86_64 0:3.7.16-1.amzn2.0.4 will be installed  
--> Processing Dependency: python3-rpm-macros for package: python3-devel-3.7.16-1.amzn2.0.4.x86_64  
--> Processing Dependency: system-rpm-config for package: python3-devel-3.7.16-1.amzn2.0.4.x86_64  
--> Running transaction check
```

Installation completed

```
root@ip-172-31-47-204:~  
Installing : python-srpm-macros-3-60.amzn2.0.1.noarch 5/8  
Installing : python-rpm-macros-3-60.amzn2.0.1.noarch 6/8  
Installing : python3-rpm-macros-3-60.amzn2.0.1.noarch 7/8  
Installing : python3-devel-3.7.16-1.amzn2.0.4.x86_64 8/8  
Verifying : python-srpm-macros-3-60.amzn2.0.1.noarch 1/8  
Verifying : system-rpm-config-9.1.0-76.amzn2.0.14.noarch 2/8  
Verifying : python-rpm-macros-3-60.amzn2.0.1.noarch 3/8  
Verifying : dwz-0.11-3.amzn2.0.3.x86_64 4/8  
Verifying : python3-rpm-macros-3-60.amzn2.0.1.noarch 5/8  
Verifying : go-srpm-macros-3.0.15-23.amzn2.0.2.noarch 6/8  
Verifying : python3-devel-3.7.16-1.amzn2.0.4.x86_64 7/8  
Verifying : perl-srpm-macros-1-8.amzn2.0.1.noarch 8/8  
  
Installed:  
python3-devel.x86_64 0:3.7.16-1.amzn2.0.4  
  
Dependency Installed:  
dwz.x86_64 0:0.11-3.amzn2.0.3 go-srpm-macros.noarch 0:3.0.15-23.amzn2.0.2  
perl-srpm-macros.noarch 0:1-8.amzn2.0.1 python-rpm-macros.noarch 0:3-60.amzn2.0.1  
python-srpm-macros.noarch 0:3-60.amzn2.0.1 python3-rpm-macros.noarch 0:3-60.amzn2.0.1  
system-rpm-config.noarch 0:9.1.0-76.amzn2.0.14  
  
Complete!  
[root@ip-172-31-47-204 ~]#
```

## Starting the ThriftServer: hbase thrift start

```
[root@ip-172-31-47-204 ~]# hbase thrift start  
Exception in thread "main" java.net.BindException: Port in use: 0.0.0.0:9095  
    at org.apache.hadoop.hbase.http.HttpServer.openListeners(HttpServer.java:1117)  
    at org.apache.hadoop.hbase.http.HttpServer.start(HttpServer.java:1052)  
    at org.apache.hadoop.hbase.http.InfoServer.start(InfoServer.java:100)  
    at org.apache.hadoop.hbase.thrift.ThriftServer.doMain(ThriftServer.java:104)  
    at org.apache.hadoop.hbase.thrift.ThriftServer.main(ThriftServer.java:240)  
Caused by: java.net.BindException: Address already in use  
    at sun.nio.ch.Net.bind0(Native Method)  
    at sun.nio.ch.Net.bind(Net.java:461)  
    at sun.nio.ch.Net.bind(Net.java:453)  
    at sun.nio.ch.ServerSocketChannelImpl.bind(ServerSocketChannelImpl.java:222)  
    at sun.nio.ch.ServerSocketAdaptor.bind(ServerSocketAdaptor.java:85)  
    at org.mortbay.jetty.nio.SelectChannelConnector.open(SelectChannelConnector.java:216)  
    at org.apache.hadoop.hbase.http.HttpServer.openListeners(HttpServer.java:1111)  
    ... 4 more  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]# jps  
29569 RESTServer  
30274 DataNode  
16389 Main  
29733 ThriftServer  
16390 Main  
22057 Jps  
7500 MRAppMaster
```

## Installing Happy base & Cython (Pre-requisite for Happy Base):

pip install Cython

pip install happybase

```
root@ip-172-31-47-204:~  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]# pip install Cython  
WARNING: Running pip install with root privileges is generally not a good idea. Try 'pip3 install --user' instead.  
Collecting Cython  
  Downloading https://files.pythonhosted.org/packages/e3/7f/f584f5d15323feb897d42ef0e9d910649e2150d7a30cf7e7a8ccld236e6f/Cython-3.0.8-py2.py3-none-any.whl (1.2MB)  
    100% |#####| 1.2MB 834KB/s  
Installing collected packages: Cython  
Successfully installed Cython-3.0.8  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]# pip install happybase  
WARNING: Running pip install with root privileges is generally not a good idea. Try 'pip3 install --user' instead.  
Collecting happybase  
  Using cached https://files.pythonhosted.org/packages/d1/9c/f5f7bdb5439cda2b7da4e20ac24ec0e2455fd68aade8397f211d2994c39d/happybase-1.2.0.tar.gz  
Requirement already satisfied: six in /usr/local/lib/python3.7/site-packages (from happybase)  
Collecting thriftpy2>=0.4 (from happybase)  
  Using cached https://files.pythonhosted.org/packages/44/c3/20664039450f04a5630b68daaa00d539c9cd5338a17d5a28c3a553c10de2/thriftpy2-0.4.20.tar.gz  
Collecting ply<4.0,>=3.4 (from thriftpy2>=0.4->happybase)  
  Downloading https://files.pythonhosted.org/packages/a3/58/35da89ee790598a0700ea49b2a66594140f44dec458c07e8e3d4979137fc/ply-3.11-py2.py3-none-any.whl (49kB)  
    100% |#####| 51kB 7.0MB/s  
Installing collected packages: ply, thriftpy2, happybase  
  Running setup.py install for thriftpy2 ... done  
  Running setup.py install for happybase ... done  
Successfully installed happybase-1.2.0 ply-3.11 thriftpy2-0.4.20  
[root@ip-172-31-47-204 ~]#
```

Developed the Python Script, Bulk\_Import.py (later renamed to Batch\_Ingest.py to comply with the file naming convention)

Executed the MR Script for bulk import of 2 files in the dataset to HBase table using following command:



python Bulk\_Import.py

```
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]#  
[root@ip-172-31-47-204 ~]# python Bulk_Import.py  
starting batch insert of yellow_tripdata_2017-03.csv
```

Checked the statistics of the HBase table after bulk import:

```
hbase(main):001:0>  
hbase(main):002:0*  
hbase(main):003:0* list  
TABLE  
yellow_tripdata_hbase  
1 row(s) in 0.2720 seconds  
  
=> ["yellow_tripdata_hbase"]  
hbase(main):004:0> count 'yellow_tripdata_hbase'  
0 row(s) in 0.1540 seconds  
  
=> 0  
hbase(main):005:0> count 'yellow_tripdata_hbase'  
Current count: 1000, row: 2017-03-01 00:06:082017-03-01 00:24:54  
Current count: 2000, row: 2017-03-01 00:12:292017-03-01 00:23:51  
Current count: 3000, row: 2017-03-01 00:19:102017-03-01 00:31:39  
Current count: 4000, row: 2017-03-01 00:26:092017-03-01 00:32:08  
Current count: 5000, row: 2017-03-01 00:33:362017-03-01 00:55:38
```

```
hadoop@ip-172-31-47-204:~  
Current count: 386000, row: 2017-03-02 07:37:312017-03-02 07:44:27  
Current count: 387000, row: 2017-03-02 07:40:212017-03-02 08:01:10  
Current count: 388000, row: 2017-03-02 07:43:092017-03-02 07:48:15  
Current count: 389000, row: 2017-03-02 07:45:542017-03-02 07:52:52  
Current count: 390000, row: 2017-03-02 07:48:452017-03-02 07:55:25  
Current count: 391000, row: 2017-03-02 07:51:322017-03-02 07:59:34  
Current count: 392000, row: 2017-03-02 07:54:262017-03-02 08:14:51  
Current count: 393000, row: 2017-03-02 07:57:212017-03-02 08:13:42  
Current count: 394000, row: 2017-03-02 08:00:302017-03-02 08:28:20  
Current count: 395000, row: 2017-03-02 08:04:022017-03-02 08:13:35  
Current count: 396000, row: 2017-03-02 08:07:212017-03-02 08:15:24  
Current count: 397000, row: 2017-03-02 08:10:402017-03-02 08:12:38  
Current count: 398000, row: 2017-03-02 08:14:132017-03-02 08:21:35  
Current count: 399000, row: 2017-03-02 08:17:442017-03-02 08:34:58  
Current count: 400000, row: 2017-03-02 08:21:182017-03-02 08:27:49  
Current count: 401000, row: 2017-03-02 08:24:522017-03-02 08:27:18  
Current count: 402000, row: 2017-03-02 08:28:252017-03-02 08:35:35  
Current count: 403000, row: 2017-03-02 08:32:022017-03-02 08:42:15  
Current count: 404000, row: 2017-03-02 08:35:352017-03-02 08:41:27  
Current count: 405000, row: 2017-03-02 08:39:082017-03-02 08:46:15  
Current count: 406000, row: 2017-03-02 08:42:502017-03-02 08:59:41  
Current count: 407000, row: 2017-03-02 08:46:202017-03-02 08:56:55  
Current count: 408000, row: 2017-03-02 08:49:472017-03-02 09:16:38  
Current count: 409000, row: 2017-03-02 08:53:212017-03-02 09:58:03  
Current count: 410000, row: 2017-03-02 08:57:062017-03-02 09:59:21  
410774 row(s) in 37.6070 seconds  
  
=> 410774  
hbase(main):006:0>
```

```
[hadoop@ip-172-31-47-204 ~]$ ls -lrt  
total 3611176  
-rw-rw-r-- 1 hadoop hadoop 4079310 Aug 7 2021 mysql-connector-java-8.0.25.tar.gz  
-rw-rw-r-- 1 hadoop hadoop 914029540 Nov 25 2022 yellow_tripdata_2017-01.csv  
-rw-rw-r-- 1 hadoop hadoop 863487050 Nov 25 2022 yellow_tripdata_2017-02.csv  
-rw-rw-r-- 1 hadoop hadoop 969809025 Nov 25 2022 yellow_tripdata_2017-03.csv  
-rw-rw-r-- 1 hadoop hadoop 946349441 Nov 25 2022 yellow_tripdata_2017-04.csv  
-rwxrwx--- 1 hadoop hadoop 1022 Feb 28 15:32 mrtask_a.py  
-rwxrwx--- 1 hadoop hadoop 1038 Feb 28 15:32 mrtask_b.py  
-rwxrwx--- 1 hadoop hadoop 1389 Feb 28 15:33 mrtask_c.py  
-rwxrwx--- 1 hadoop hadoop 1834 Feb 28 15:33 mrtask_d.py  
-rwxrwx--- 1 hadoop hadoop 1360 Feb 28 15:33 mrtask_e.py  
-rwxrwx--- 1 hadoop hadoop 1584 Feb 28 15:33 mrtask_f.py  
drwxr-xr-x 3 hadoop hadoop 149 Feb 29 09:14 mysql-connector-java-8.0.25  
-rw-rw-r-- 1 hadoop hadoop 51862 Feb 29 12:12 yellow_tripdata.java  
-rwxrwx--- 1 hadoop hadoop 1820 Feb 29 15:56 Bulk_Import.py  
[hadoop@ip-172-31-47-204 ~]$ vi Bulk_Import.py
```